


What's new in Swarm 1.1

the Docker-native clustering system







Docker San Mateo meetup - 02/17/2016

@vieux


Running containers on multiple hosts



Without Docker Swarm



With Docker Swarm



Swarm in a nutshell

- Exposes several Docker Engines as a single virtual Engine
- Serves the standard Docker API
- Extremely easy to get started
- Batteries included but swappable



Setup: Without networking



Setup using the hosted discovery service

/!\ Not to be used in production, for testing only /!\

- Create a cluster:
`$ swarm create`
- Add nodes to a cluster:
`$ swarm join --advertise=<engine_ip>:<engine_port> token://<token>`
- Start Swarm
`$ swarm manage <...> token://<token>`



Setup using your own K/V store

- Add nodes to a cluster:
`$ swarm join --advertise=<engine_ip>:<engine_port> \`
`consul://<ip_consul>:<port_consul>`
- Start Swarm
`$ swarm manage <...> consul://<ip_consul>:<port_consul>`

You can also use etcd or zookeeper



Setup using a file (static list of nodes)

- Add nodes to the file:
\$ echo 10.0.0.1:2375 > my_cluster
\$ echo 10.0.0.2:2375 >> my_cluster
\$ echo 10.0.0.3:2375 >> my_cluster
- Start Swarm
\$ swarm manage <...> file:///my_cluster



Setup: With networking



Setup using your own K/V store

- Configure networking on engine
\$ **docker daemon --cluster-advertise=<engine_ip>:<engine_port> **
\$ **--cluster-store=consul://<ip_consul>:<port_consul>**
- Start Swarm
\$ **swarm manage <...> --discovery-opt kv.path=docker/docker **
\$ **consul://<ip_consul>:<port_consul>**

You can also use etcd or zookeeper



Docker Swarm internals



Resource Management

- Memory
\$ **docker run -m 1g ...**
- CPU
\$ **docker run -c 1 ...**
- Ports
\$ **docker run -p 80:80 ...**



Constraints

- Standard constraints induced from docker info
docker run -e "constraint:operatingsystem==*fedora*" ...
docker run -e "constraint:storagedriver==*aufs*" ...
- Custom constraints with host labels
docker daemon --label "region=us-east"
docker run -e "constraint:region==us-east" ...
- Pin a container to a specific host
docker run -e "constraint:node==node-2" ...



Affinities

- Containers affinities
docker run --name web nginx
docker run -e “affinity:container==web” logger
- Containers Anti-affinities
docker run --name redis-master redis
docker run --name redis-slave -e “affinity:container!=redis*” ...
- Images affinities
docker run -e “affinity:image==redis” redis



Soft Affinities/Constraints

- Soft Constraints
 - `docker run -e "constraint:operatingsystem==~*fedora*" ...`
 - `docker run -e "constraint:region==~us-east" ...`
- Soft Affinities
 - `docker run --name redis-master redis`
 - `docker run --name redis-slave -e "affinity:container!=~redis*" ...`



Swarm Scheduler

2 steps:

- 1- Apply filters to exclude nodes
 - ports
 - constraints
 - affinity
 - health
 - dependency
- 2- Use a strategy to rank and pick the best node
 - binpack
 - spread
 - random



New in Swarm 1.1

- Improved node management
- Rescheduling (EXPERIMENTAL)
- New events



Improved node management: docker info

Nodes: 3

(unknown): 10.0.0.9:2375

...

└ Status: Pending

node-1: 10.0.0.1:2375

└ Status: Unhealthy

...

└ Error: Cannot connect to the docker engine endpoint

└ UpdatedAt: 2016-02-09T19:52:56Z

node-2: 10.0.0.0:2375

└ Status: Healthy

...

└ Labels: kernelversion=4.2.0-23-generic, operatingsystem=Ubuntu 14.04.3 LTS, ...

└ Error: (none)

└ UpdatedAt: 2016-02-09T19:52:56Z

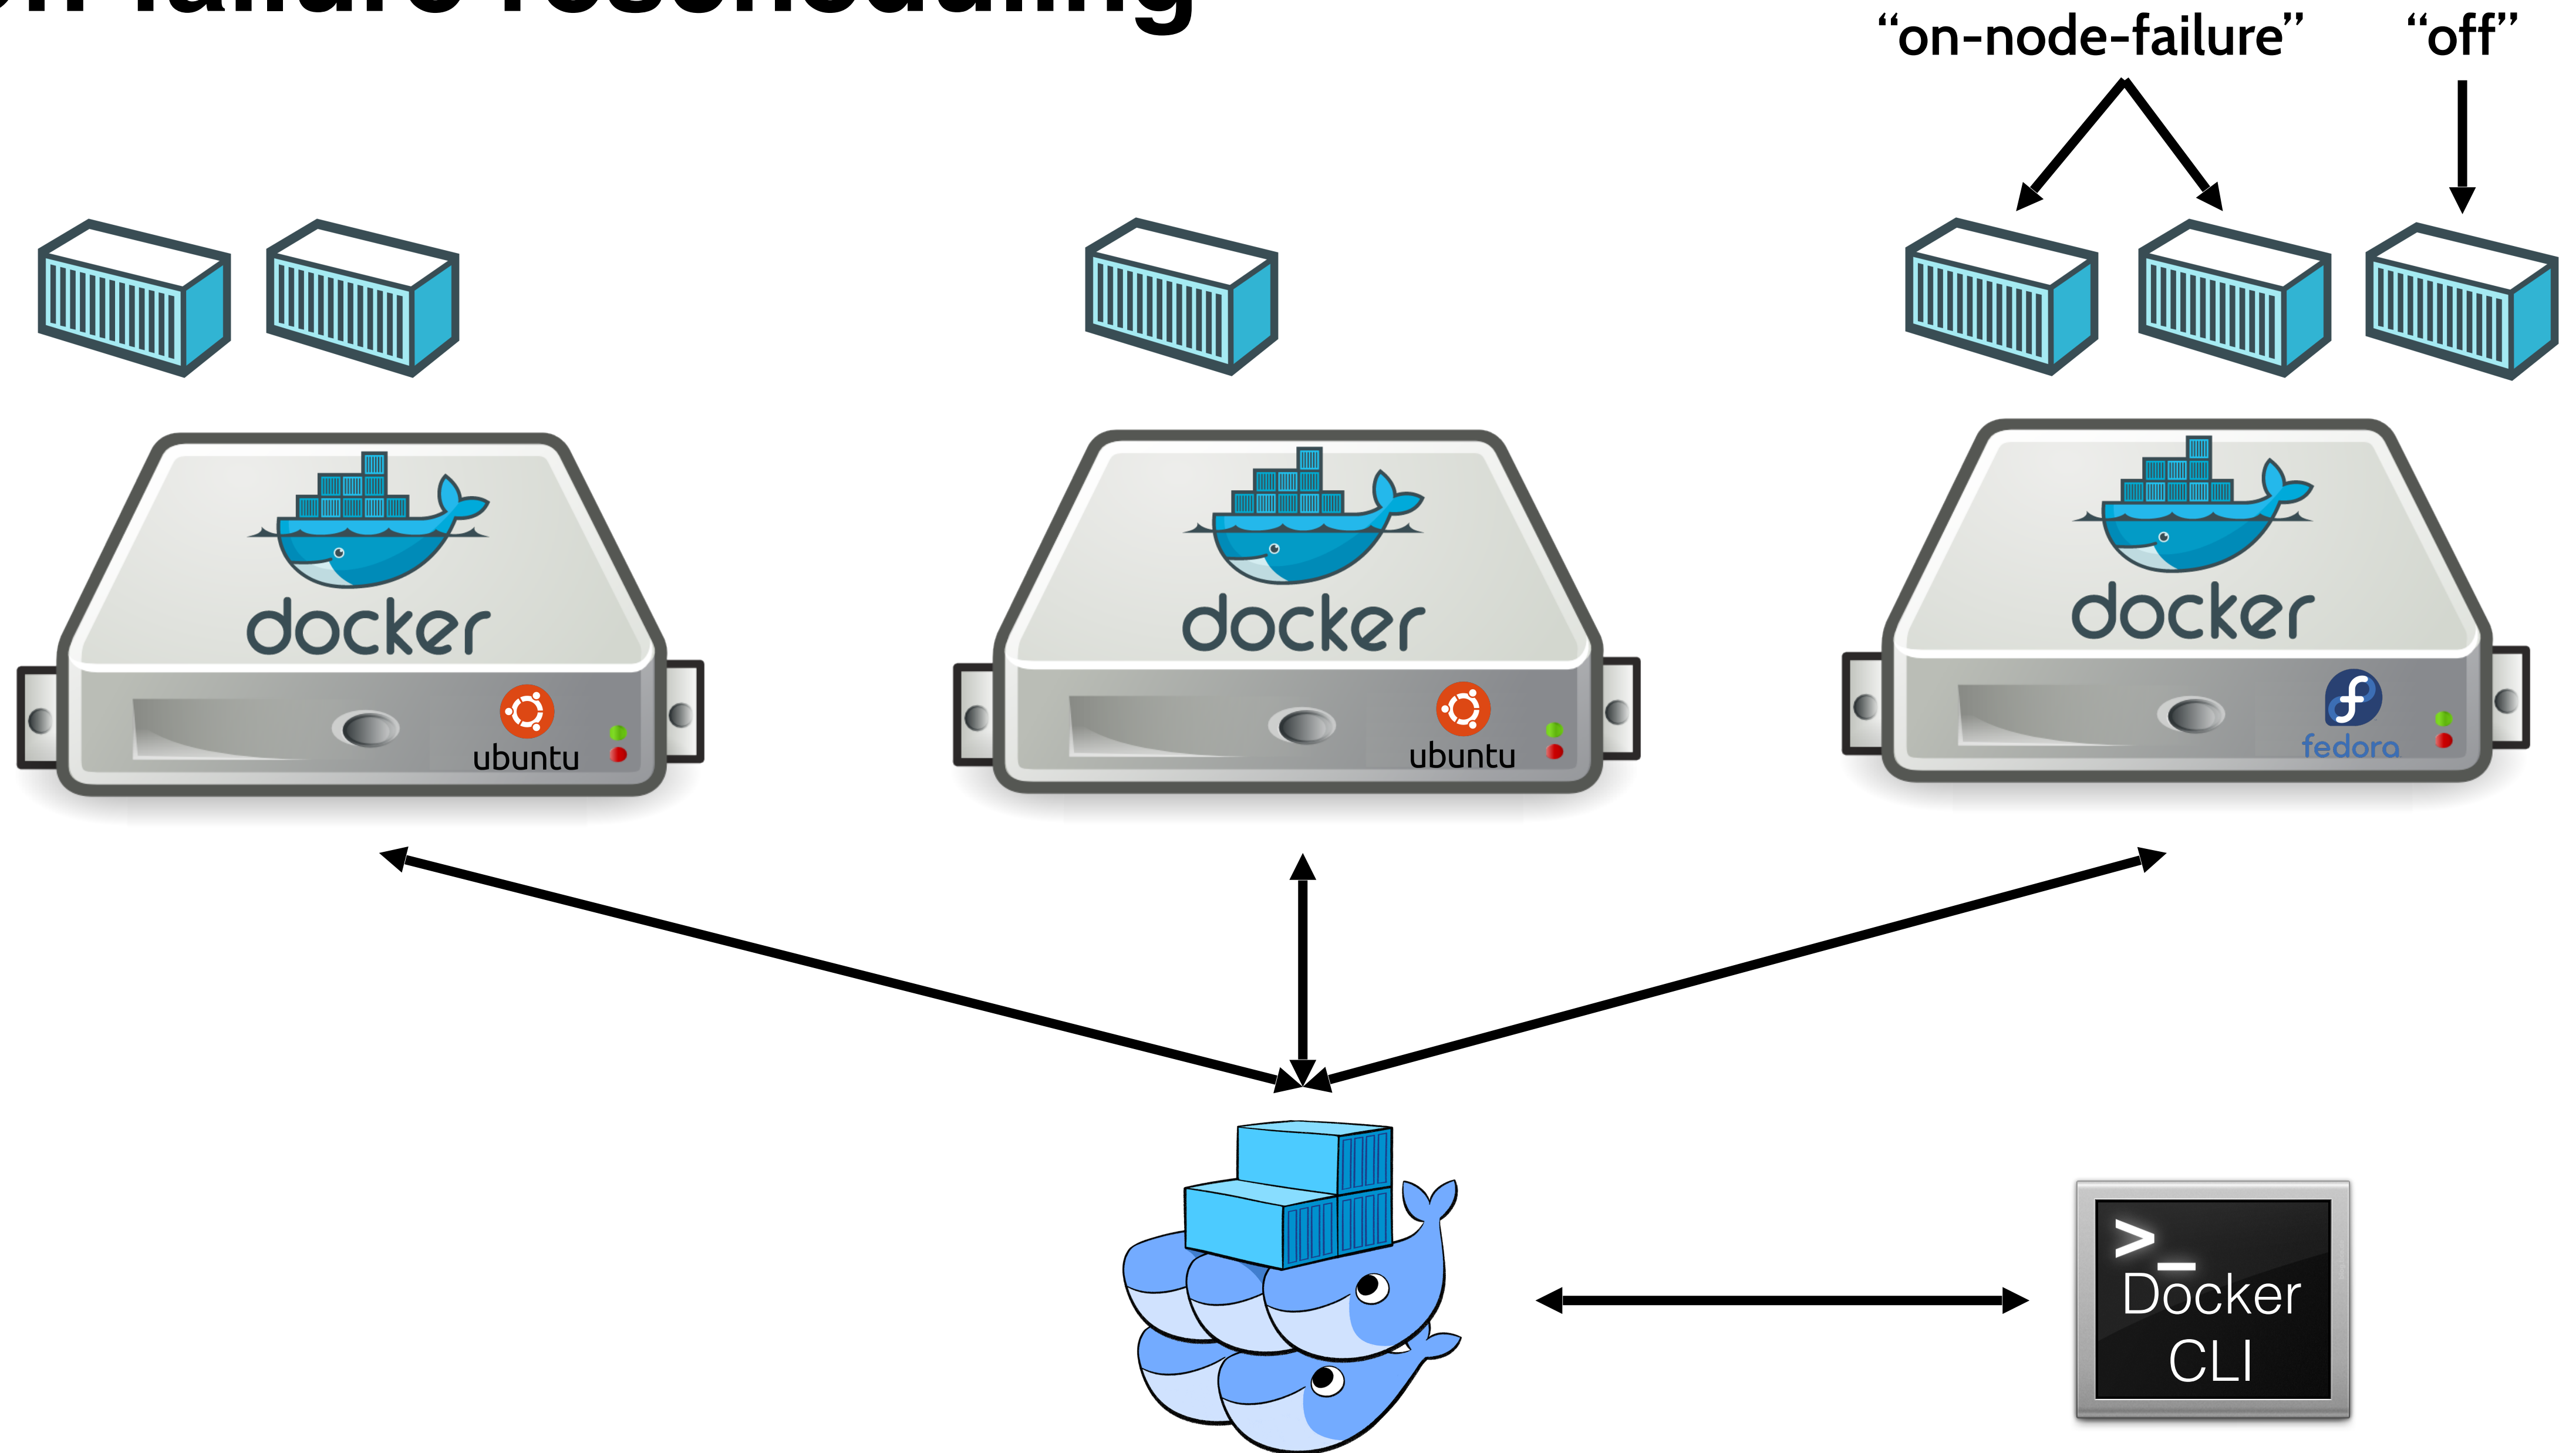


Rescheduling

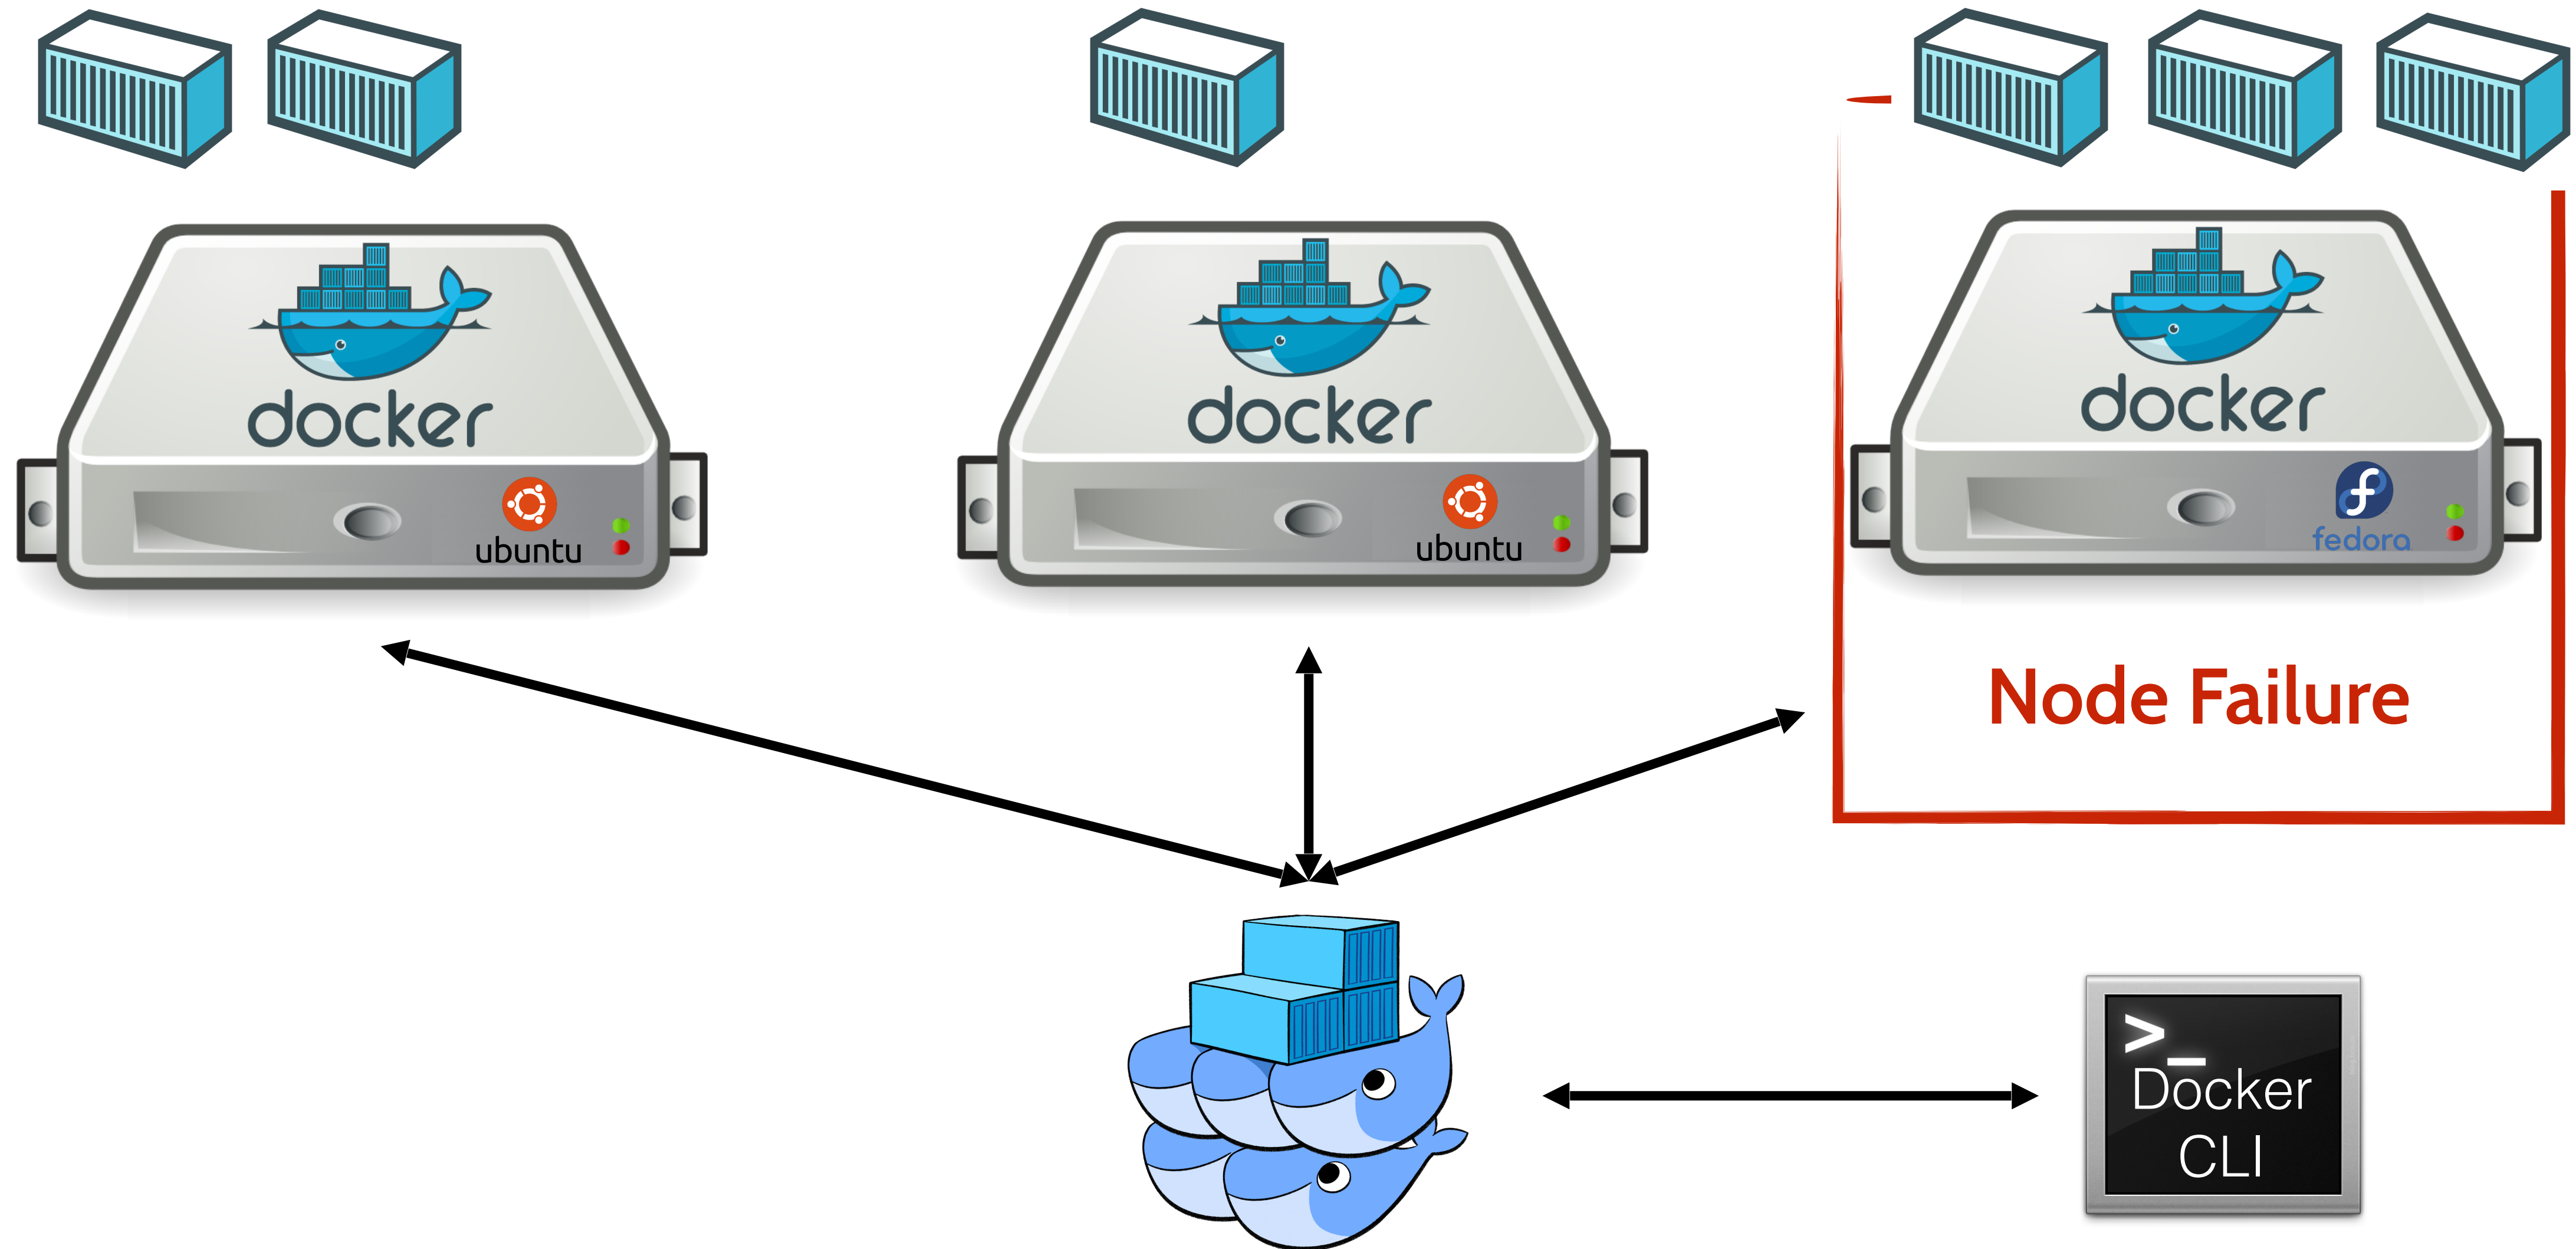
- Experimental feature
swarm --experimental manage ...
- On node failure
docker run -e "reschedule:on-node-failure" ...



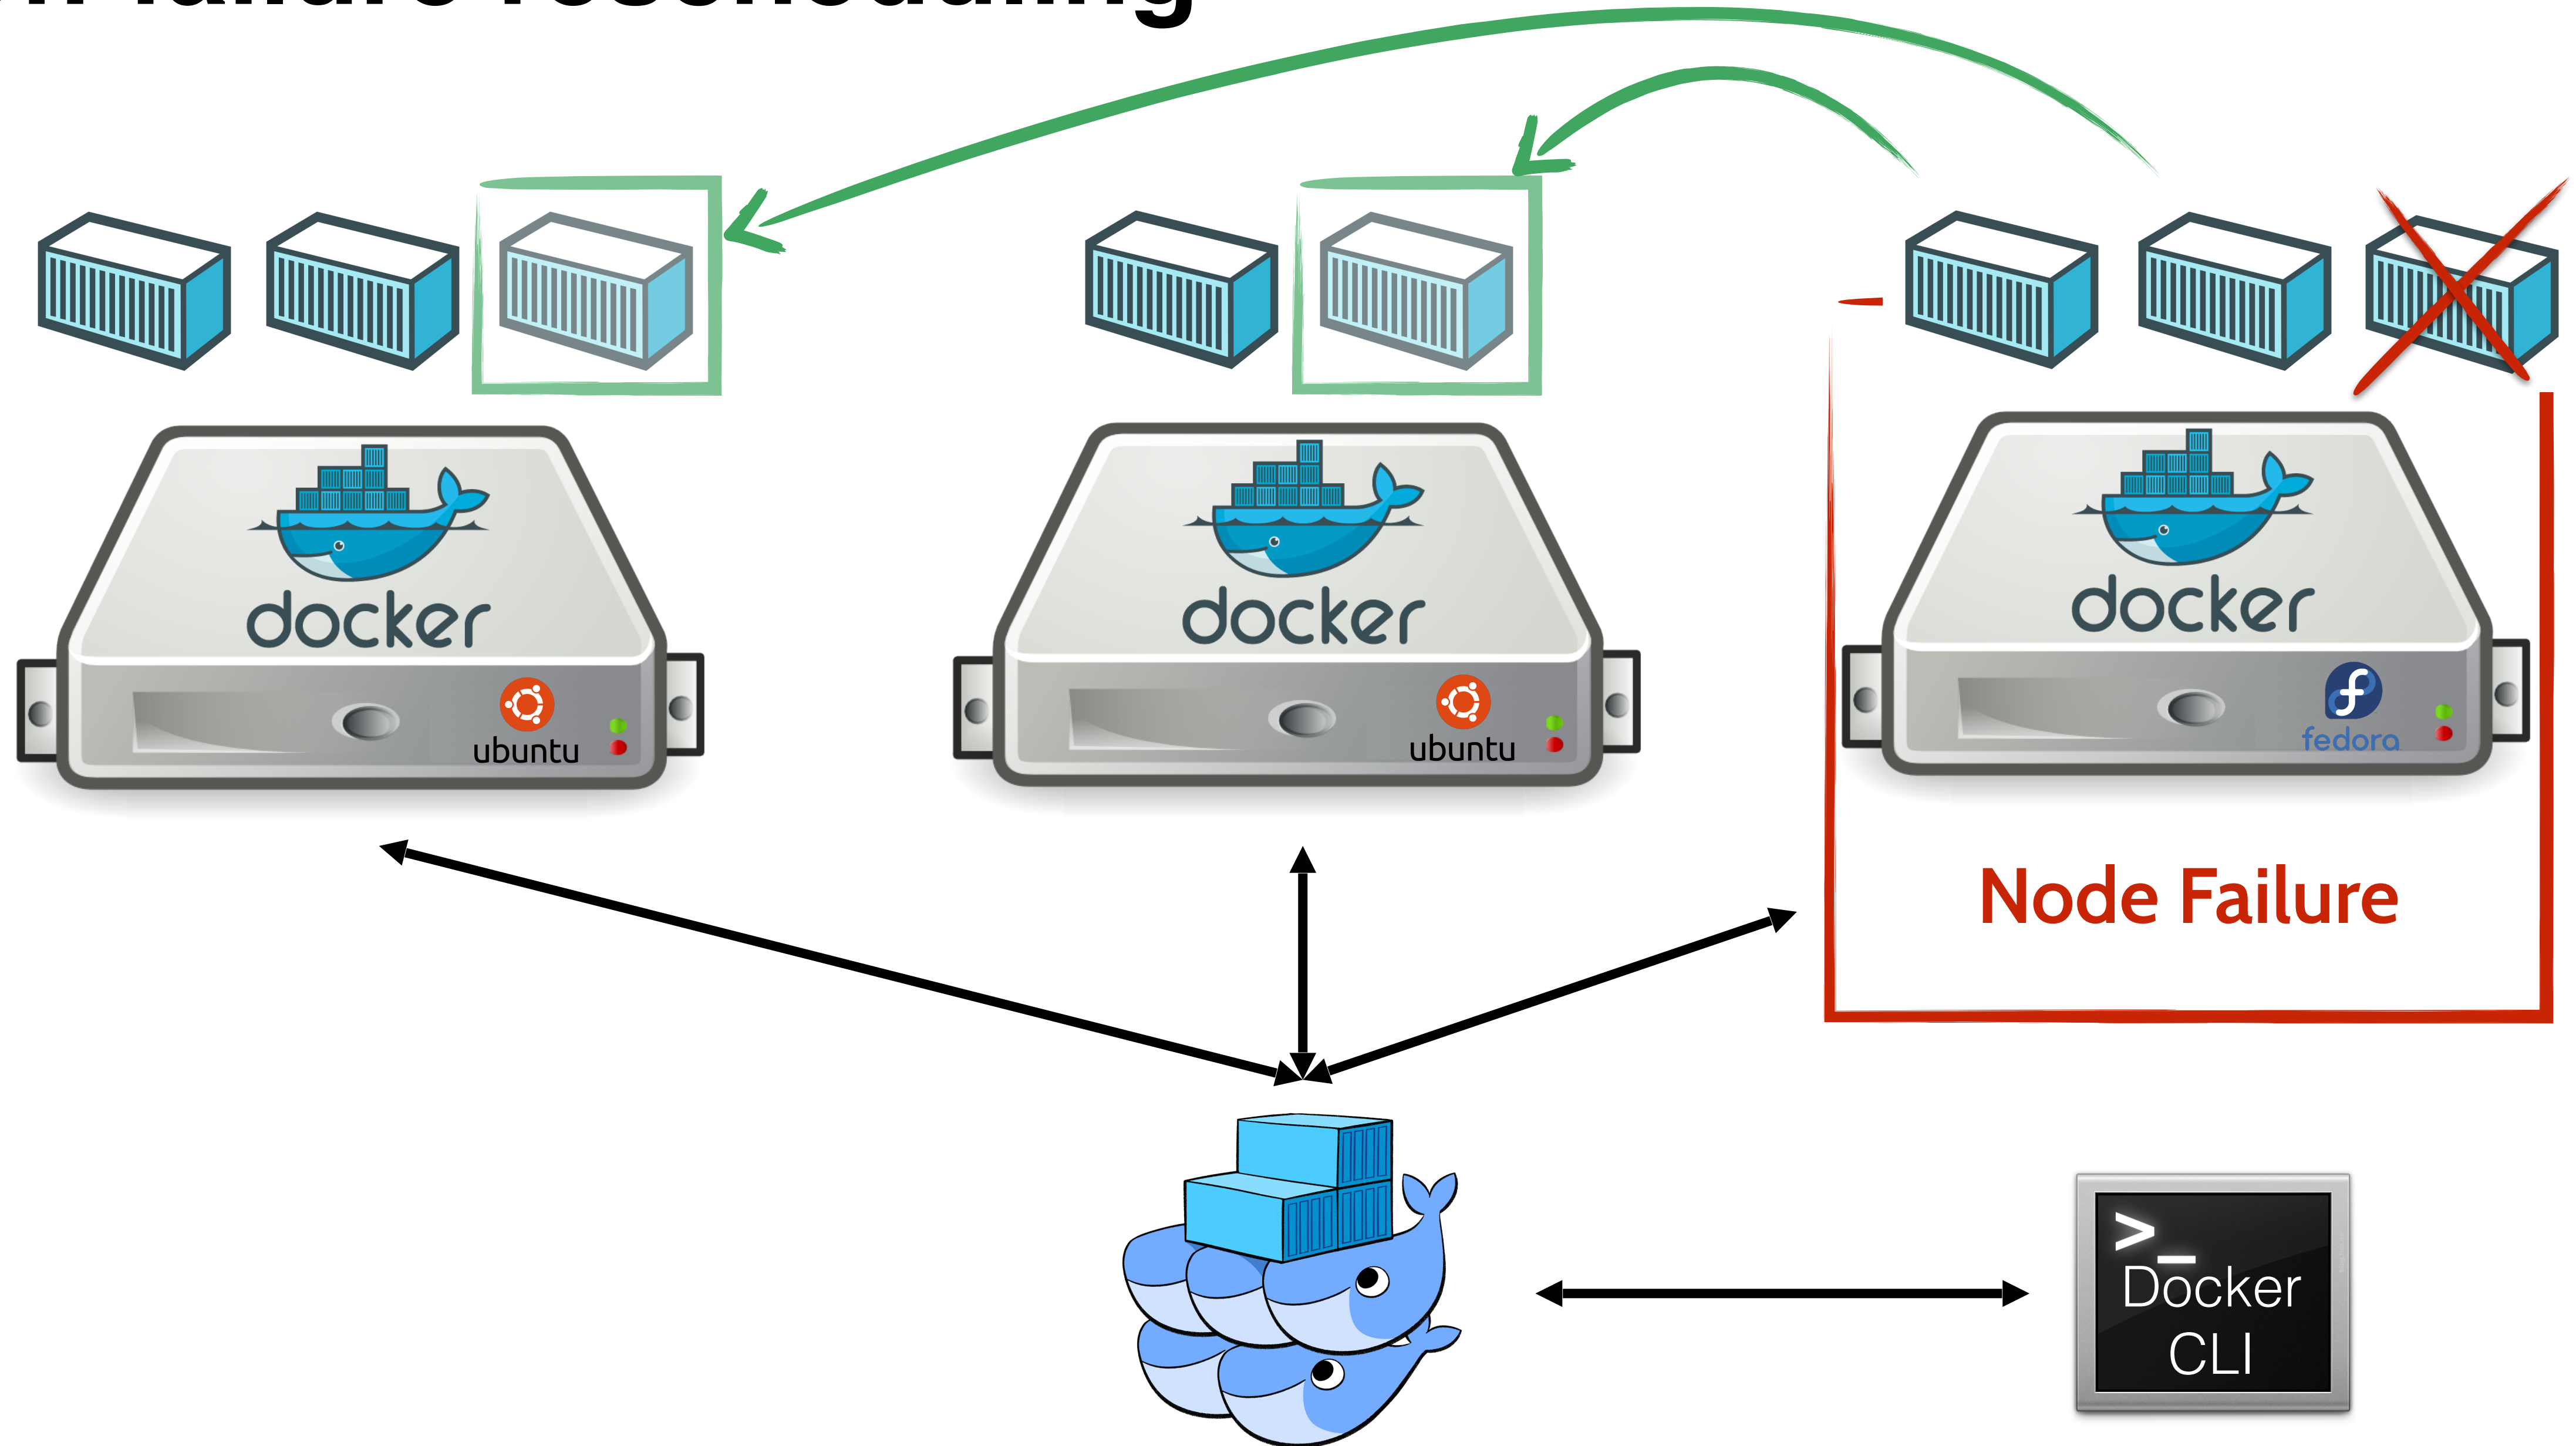
On-failure rescheduling



On-failure rescheduling



On-failure rescheduling



New events

```
$ docker events
```

```
...
```

```
2016-02-09T12:02:01 container create XXX (com.docker.swarm.id=YYY, image=busybox, node.addr=10.0.0.1:2375, node.name=node-1)
```

```
...
```

```
2016-02-09T12:02:01 network connect ZZZ (node.name=node-1, type=bridge, container=XXX, name=bridge, node.addr=10.0.0.1:2375,)
```

```
...
```

```
2016-02-09T12:03:10 swarm engine_connect (node.name=node-2, node.addr=10.0.0.2:2375)
```

```
...
```




Demo

 = RAM: 1GB
CPU: 1core

env=prod

env=stage




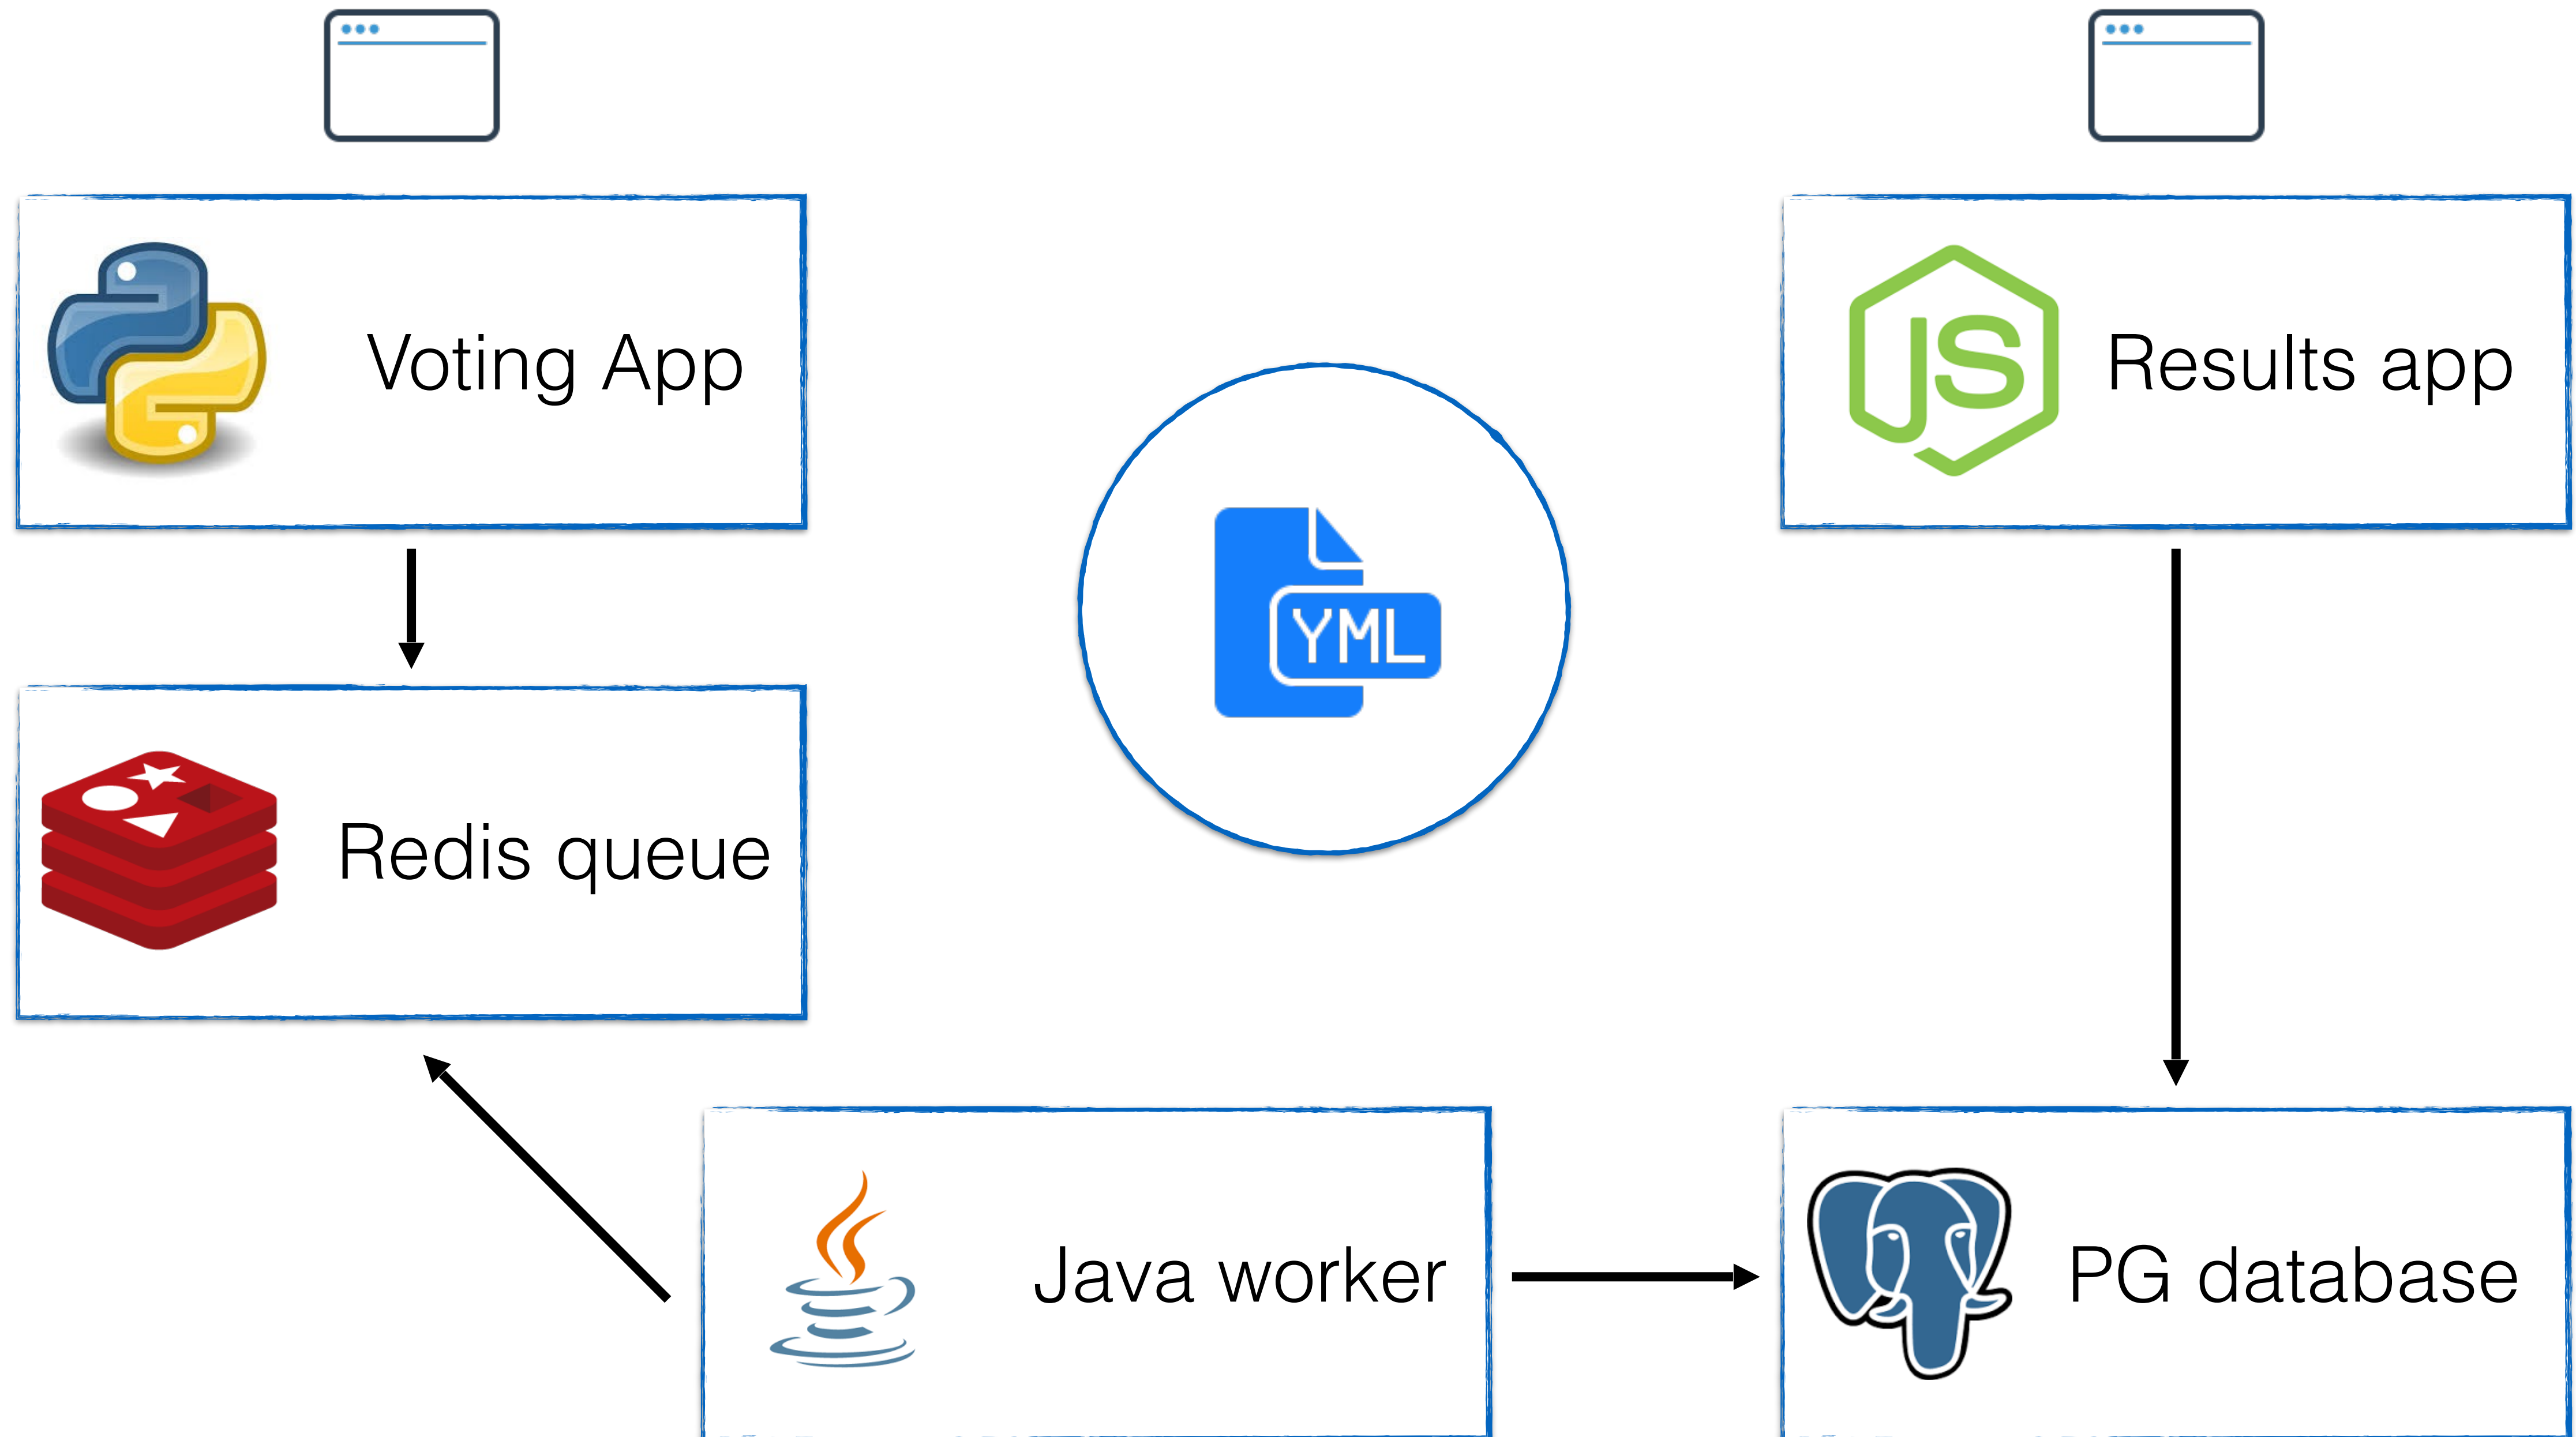
Demo

 = RAM: 1GB
CPU: 1core

env=prod

env=stage





“vote” network


10.0.1.0/24




Voting App



Redis queue



Java worker



PG database

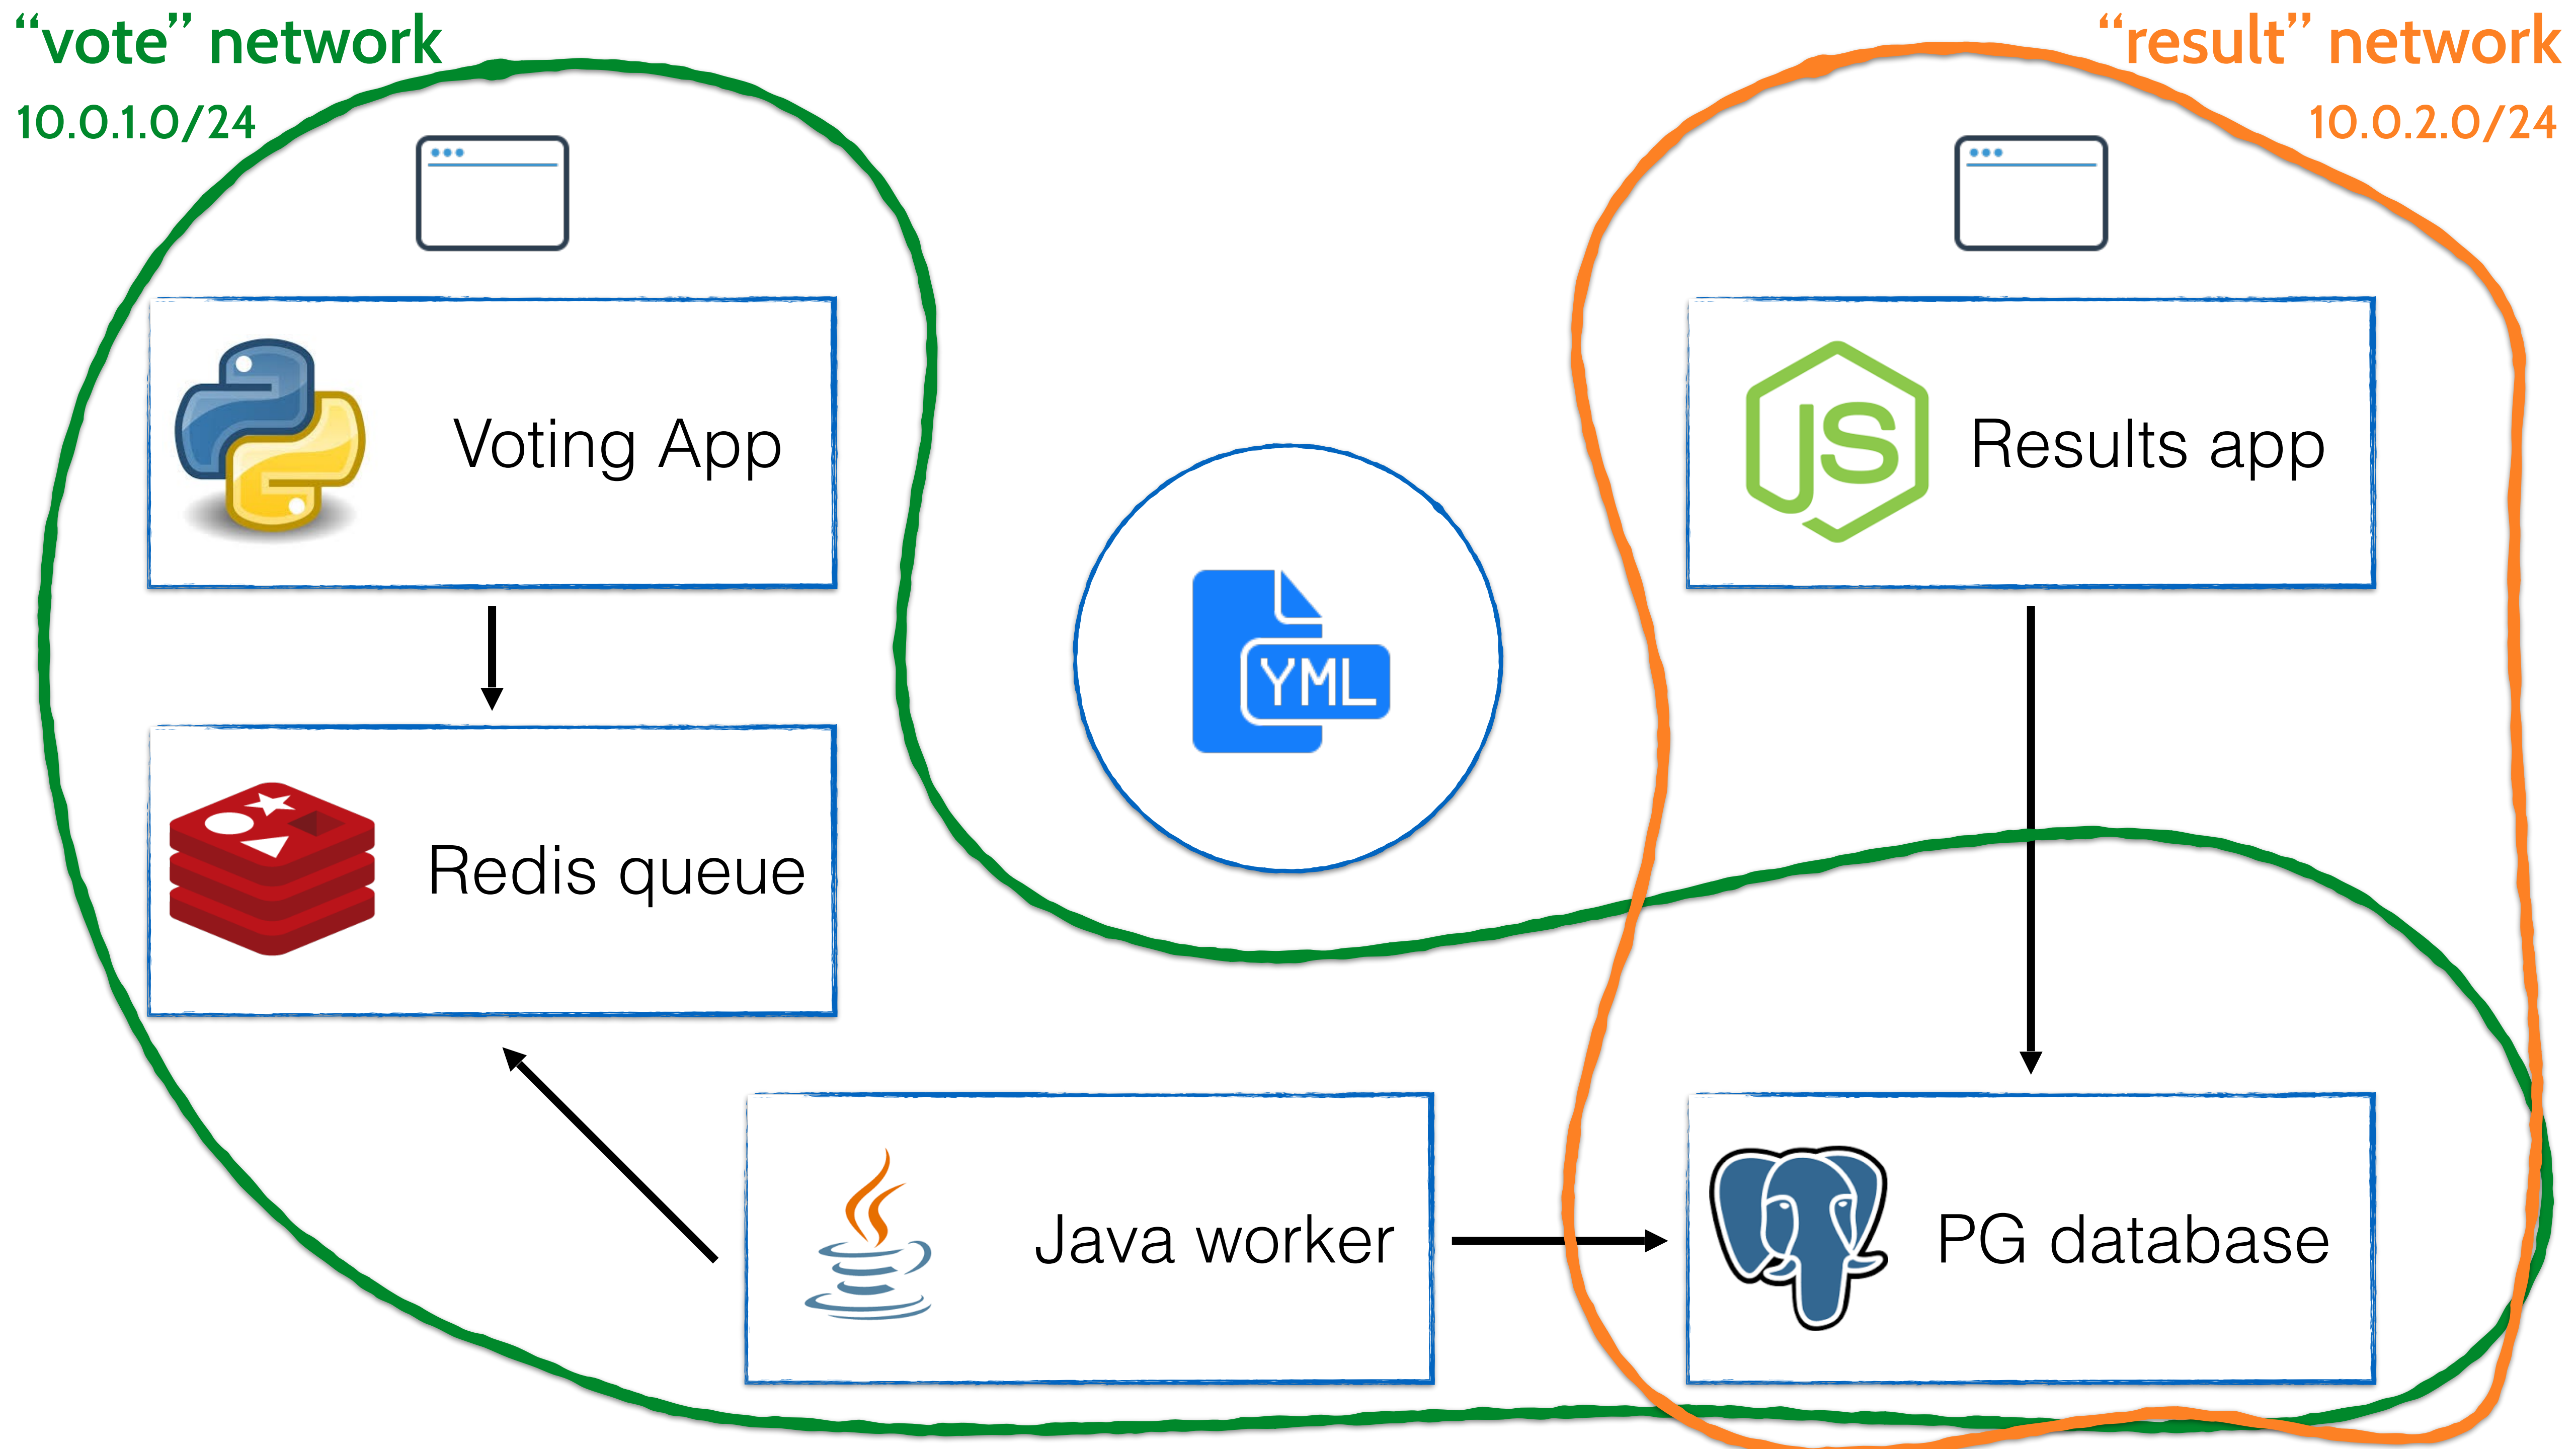


“result” network

10.0.2.0/24




Results app



Thank You. Questions?

<http://github.com/docker/swarm>

#docker-swarm on freenode



@vieux